

The Golem logo, consisting of the word "golem" in a white, lowercase, sans-serif font.

Solutions for trusted and private computations in Golem and in the wider Ecosystem

DEVCON 4

PIOTR JANIUK

We need trustworthy computations

Why do we need it?

01 —————

Consensus in a trustless, decentralized network

02 —————

Remote computation

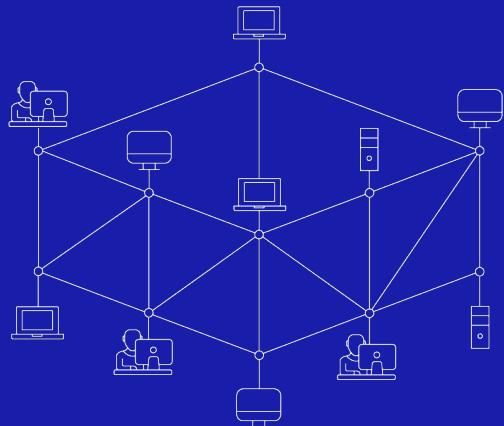
03 —————

Decentralized services

04 —————

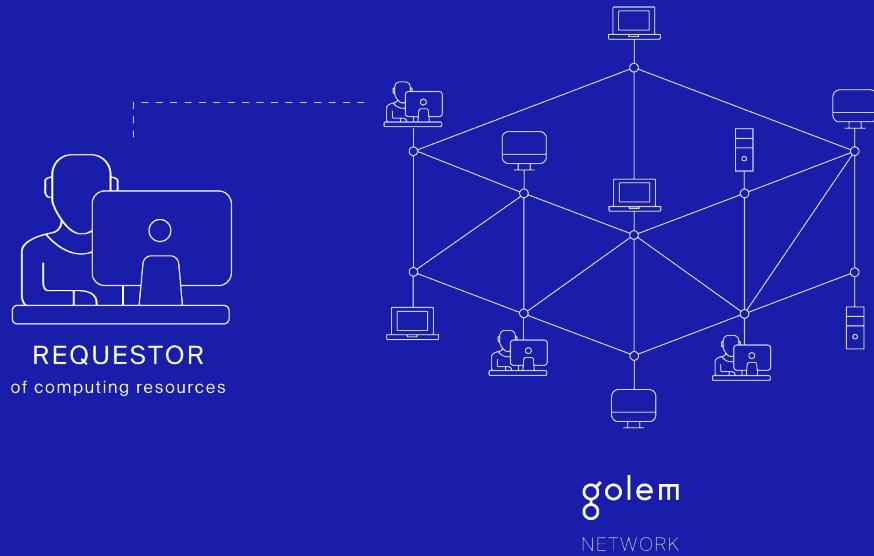
And more...

In the context of Golem

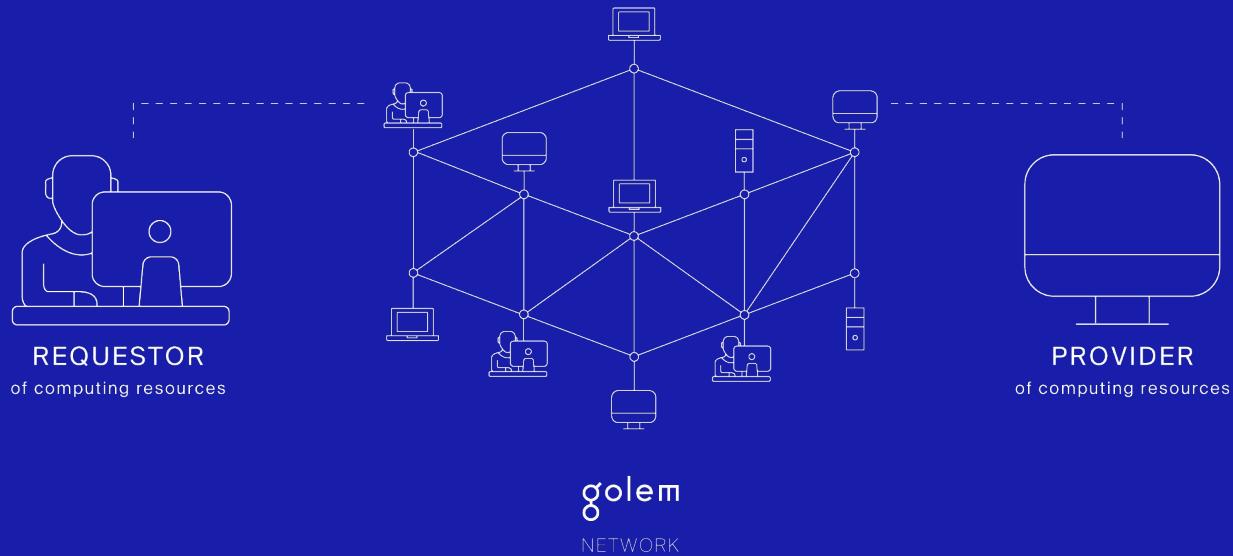


golem
NETWORK

In the context of Golem



In the context of Golem



What is Golem?

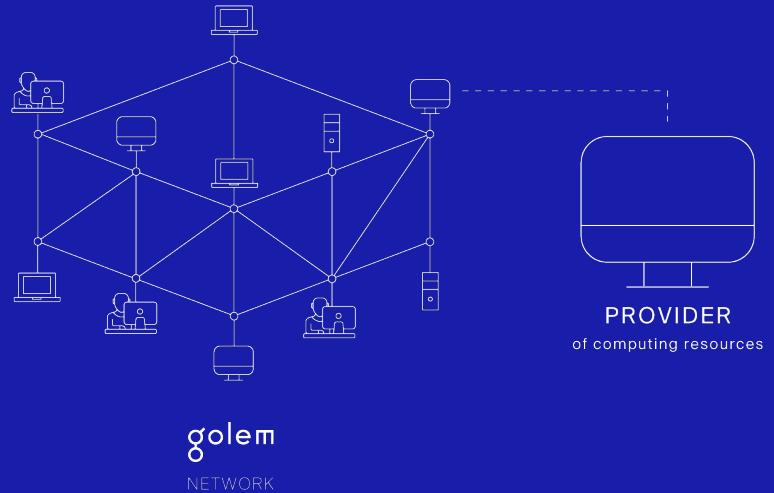
01 —————

A network of heterogeneous resources that can be either used by a **requestor** or provided to other participants by a **provider**.



02 —————

Additional layers on top of the infrastructure (e.g., economy).

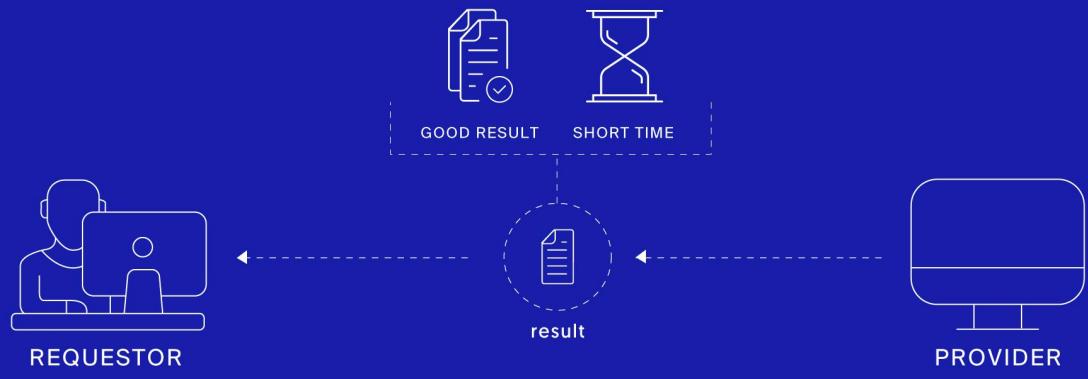


Problem statement

01 _____

Requestor wants to be sure that

- He receives valid results from the provider after a trustworthy execution on the provider's machine
- Gets results within a reasonable time

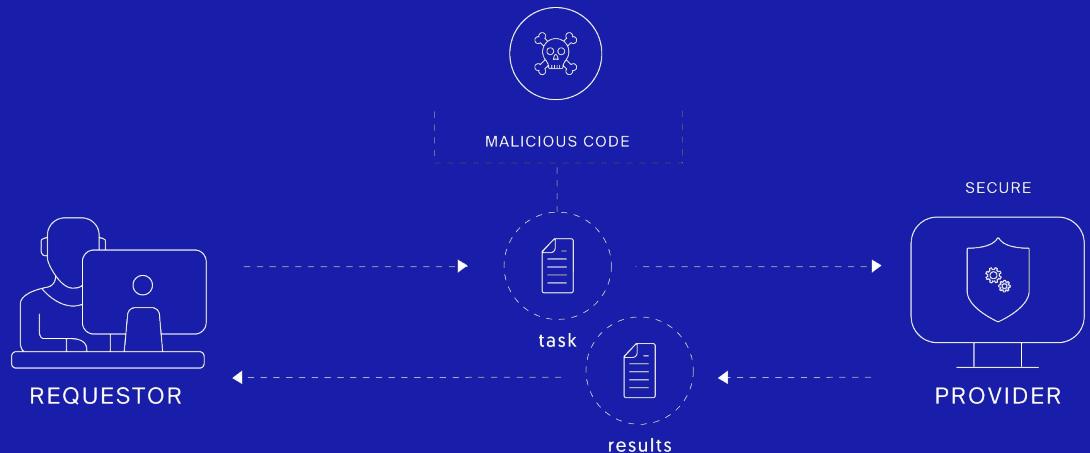


Problem statement

02

Provider on the other hand

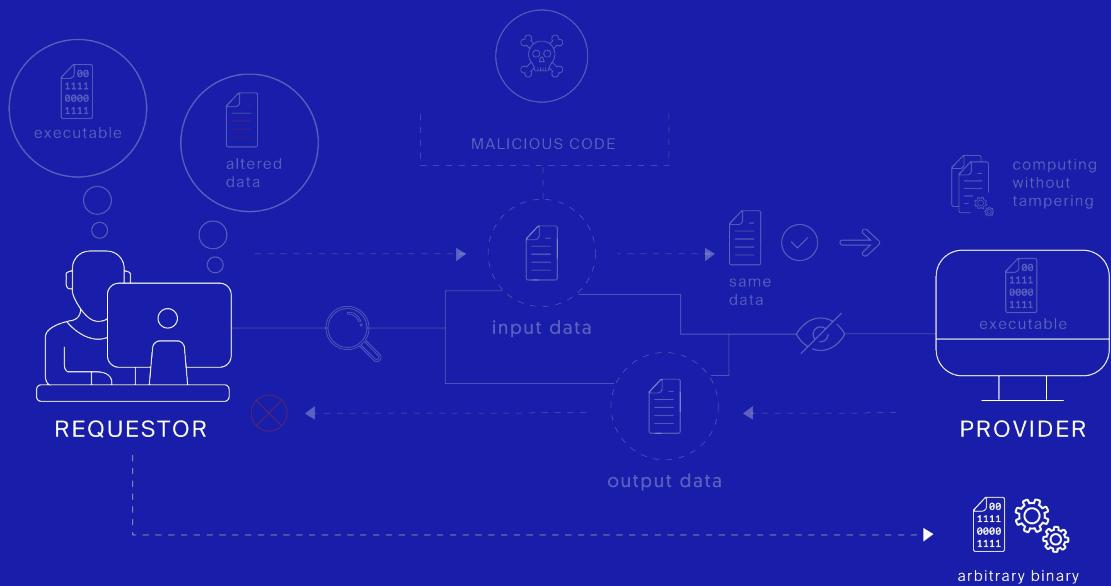
- Needs to be protected from malicious tasks



High-level requirements

01

Requestor should be able to run an arbitrary, unmodified binary efficiently

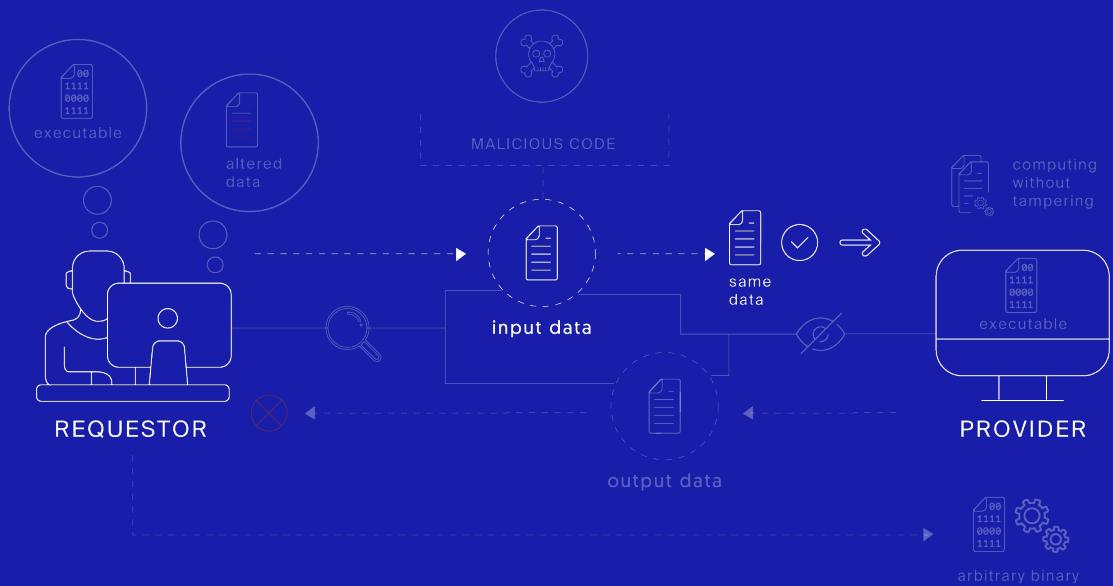


High-level requirements

02

Requestor should be able to run an arbitrary, unmodified binary efficiently

Task input data has to be the same as prepared by the requestor



High-level requirements

03

Requestor should be able to run an arbitrary, unmodified binary efficiently

Task input data has to be the same as prepared by the requestor

The binary run by a provider must be the one requestor wanted



High-level requirements

04

Requestor should be able to run an arbitrary, unmodified binary efficiently

Task input data has to be the same as prepared by the requestor

The binary run by a provider must be the one requestor wanted

The execution of the task has to be carried out without tampering



High-level requirements

05

Requestor should be able to run an arbitrary, unmodified binary efficiently

Task input data has to be the same as prepared by the requestor

The binary run by a provider must be the one requestor wanted

The execution of the task has to be carried out without tampering

Output data cannot be altered without detection



High-level requirements

06

Requestor should be able to run an arbitrary, unmodified binary efficiently

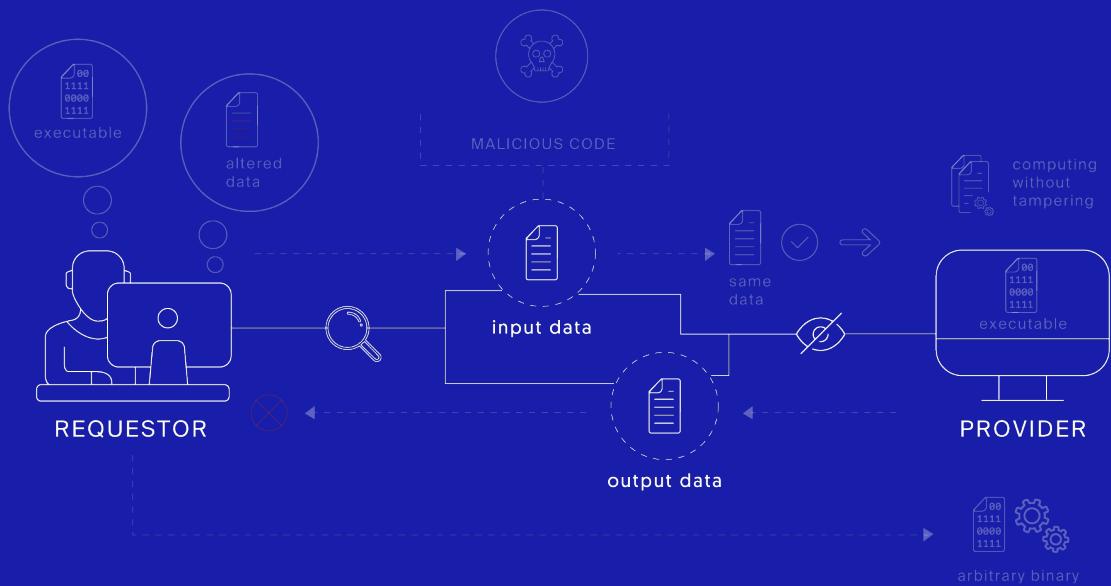
Task input data has to be the same as prepared by the requestor

The binary run by a provider must be the one requestor wanted

The execution of the task has to be carried out without tampering

Output data cannot be altered without detection

Only the requestor can see the plain text input and output data



High-level requirements

06

Requestor should be able to run an arbitrary, unmodified binary efficiently

Task input data has to be the same as prepared by the requestor

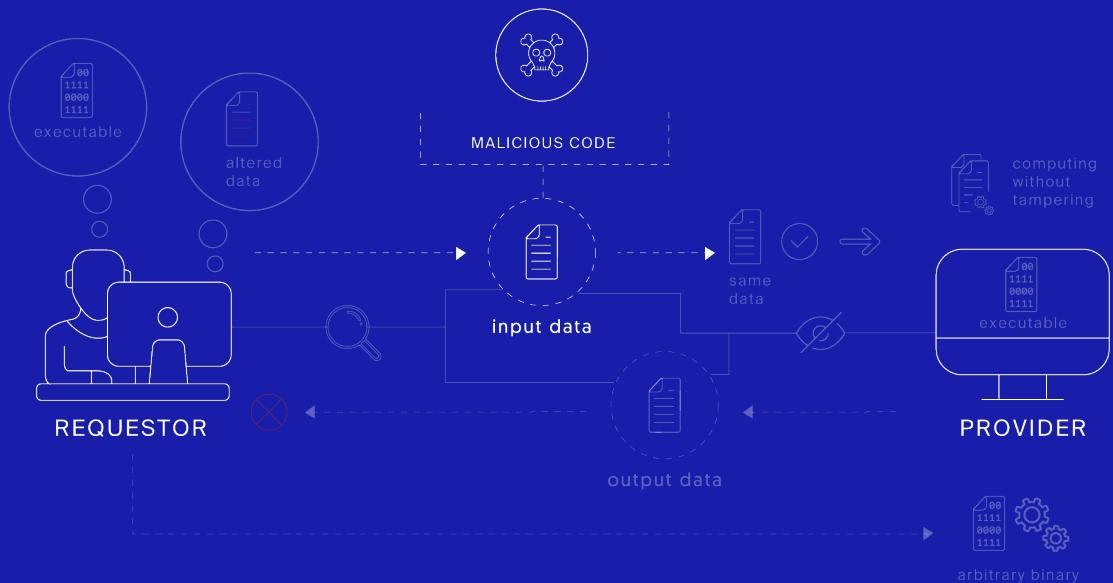
The binary run by a provider must be the one requestor wanted

The execution of the task has to be carried out without tampering

Output data cannot be altered without detection

Only the requestor can see the plain text input and output data

Provider has to be protected from malicious binaries



Meeting the requirements

01 _____

Limiting the class of problems

- deterministic (PoW or verification)
- nondeterministic (probabilistic verification)

02 _____

Introducing external sources of trust

03 _____

By means of the infrastructure (TEEs)

04 _____

Hybrid approach

Choosing the best approach

01 _____

Tradeoff

- Inherent to the method
- Task specific
- Runtime dependent

02 _____

Focus on a single one

- Based on the hardware infrastructure
- Allows remote computations

Meeting the requirements SGX, Graphene and Graphene-ng

SGX - short recap

01 —————

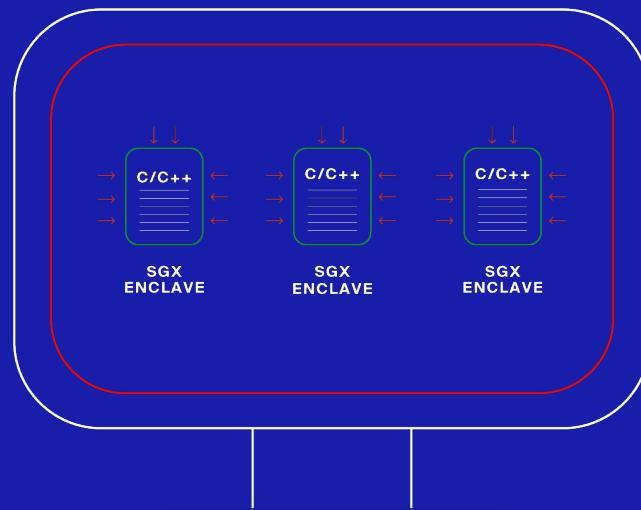
Overview

- An Intel® architecture extension designed to increase the security of application code and data
- Enclave based computation model (application is run in containers called enclaves, which are supposed to be protected from the host)
- Provides a mechanism for remote attestation

Developer perspective - an enclave

- Static interface specification
- Source code + Intel SDK
- By default limited to computations

Host (Untrusted Environment)



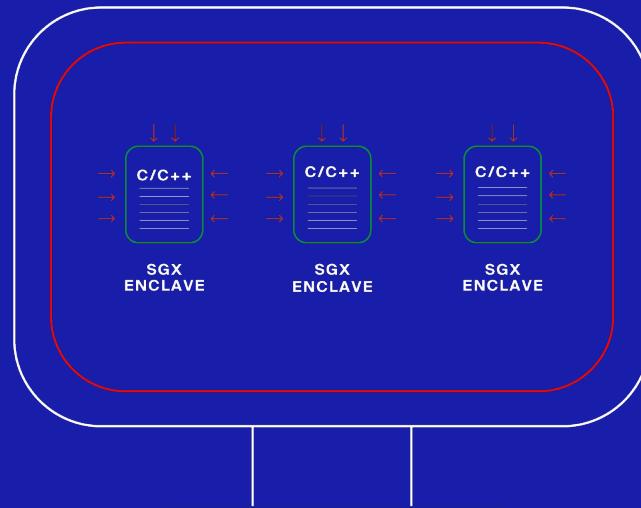
SGX - short recap

02

Summary

- Provides important security features
- Does not allow to run arbitrary applications in the enclave
- May be treated as a building block for more generic solutions

Host (Untrusted Environment)



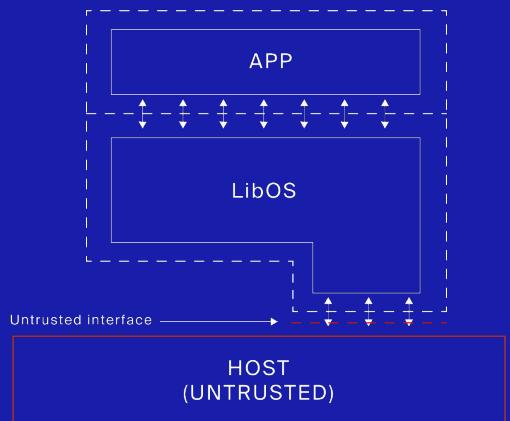
Graphene overview

01 —————

Features

- A framework with LibOS, which simulates the OS for the app in the enclave
- Runs unmodified, full Linux apps protected from the host
- Makes use of Intel SGX

Graphene architecture



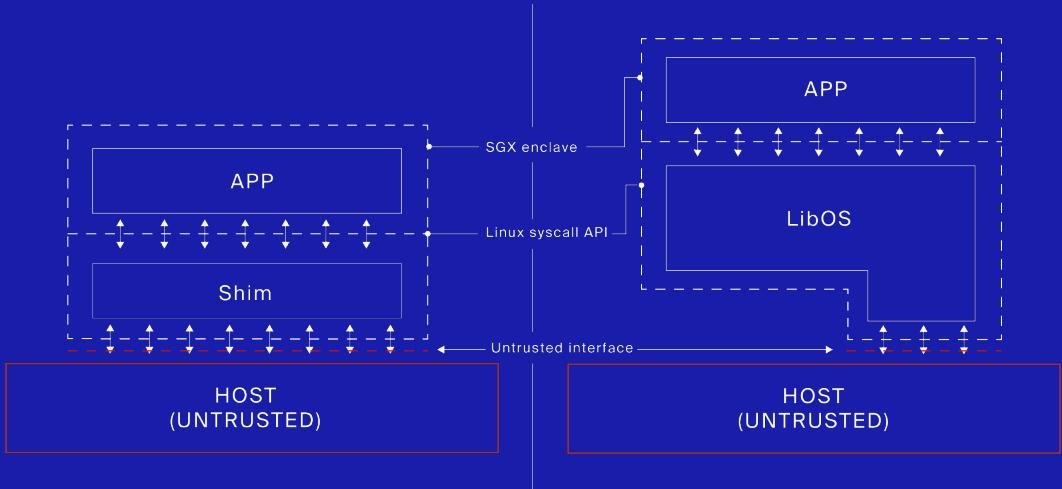
Graphene overview

02

Graphene vs other approaches

- TCB vs attack surface
- Flexibility - other host or guest OS can be supported by changing a single component
- Better suited to sandboxing

Graphene architecture

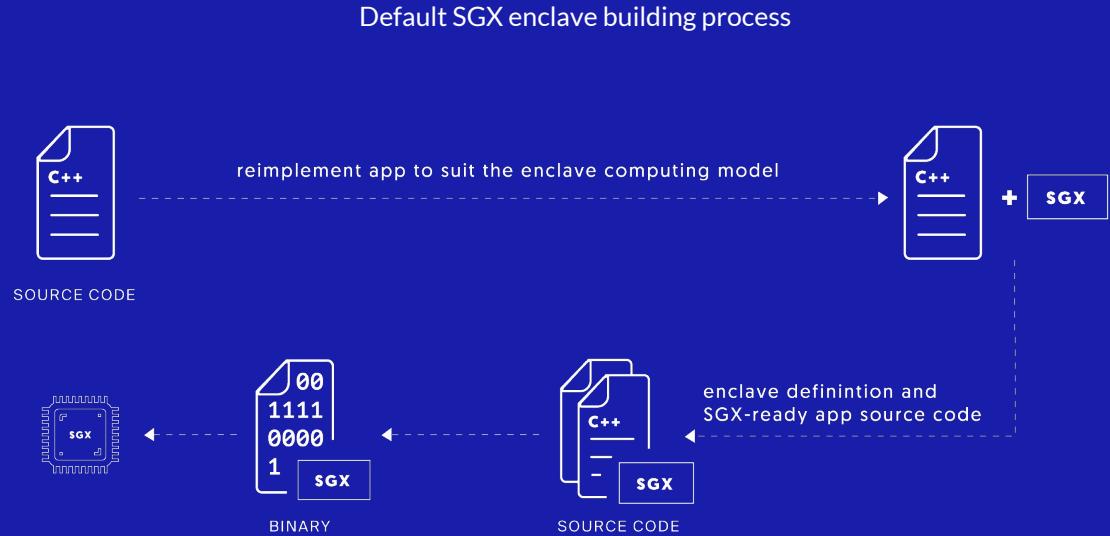


Graphene overview

03

With Graphene, a developer
doesn't have to

- Rewrite the code to suit SGX needs
- Specify the static enclave interface
- Compile the enclave

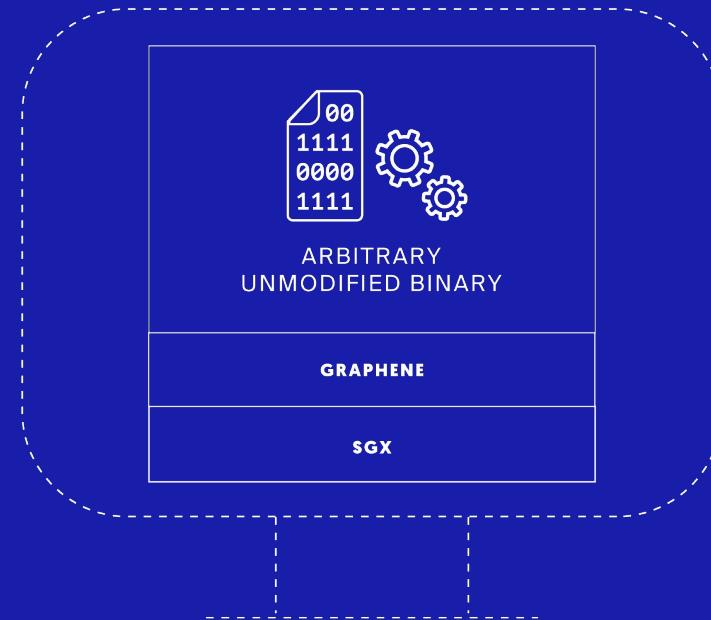


Graphene overview

04

With Graphene:

- An arbitrary, unmodified Linux binary can be run inside an enclave

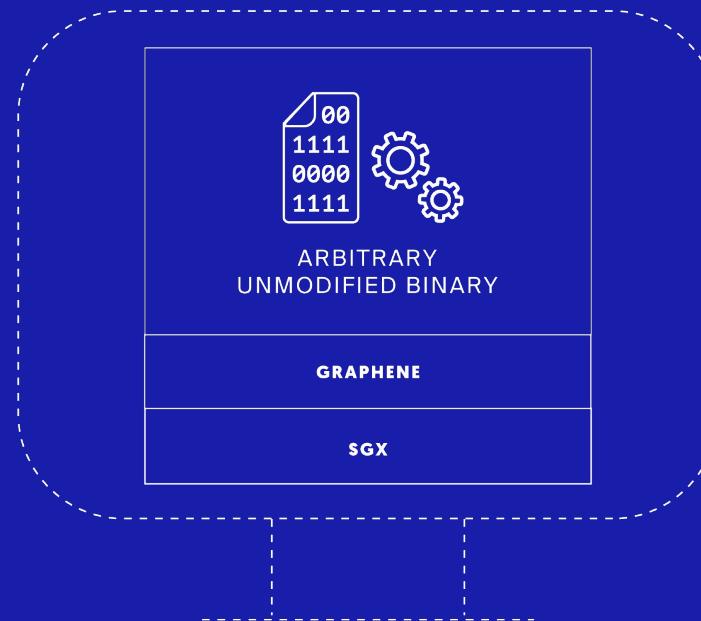


Graphene overview

05

With Graphene:

- An arbitrary, unmodified Linux binary can be run inside an enclave
- + The process still requires manual work to configure, deploy and run the application



The next step

Graphene-ng

The next step

Graphene-ng = Graphene + Protected files
+ Docker support
+ Tools (scripts)
+ Bug fixes

The next step

Graphene-ng = Graphene

Resulting in a better UX regarding interaction with the enclave for both, the developer and the end user.

- + Protected files
- + Docker support
- + Tools (scripts)
- + Bug fixes

Graphene-ng by example

PoC integration with Golem

Provider setup

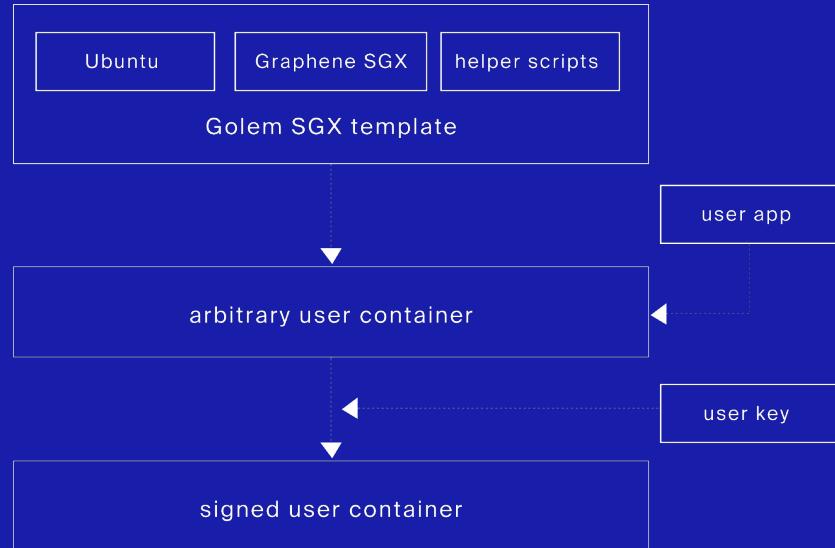
02

SGX must be enabled on provider's machine

Protected Files configuration:

- Enclave manifest
- Running the docker image

Graphene-ng Docker configuration



Provider setup

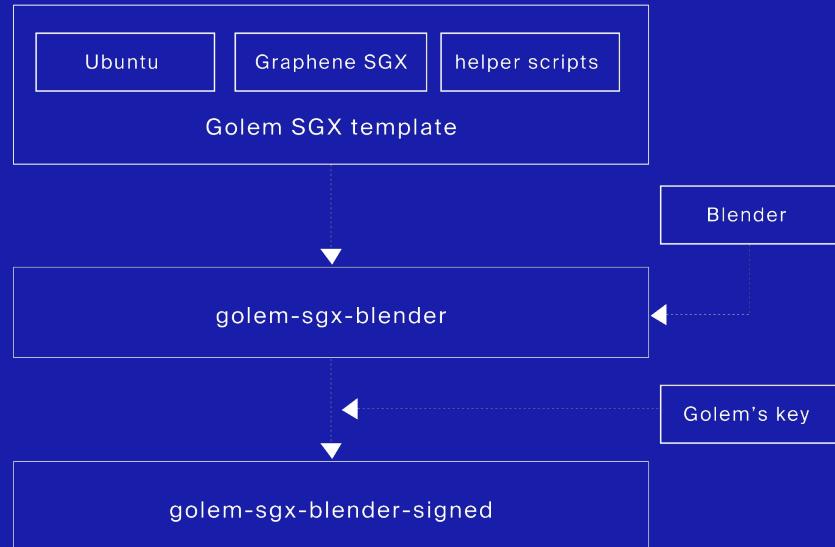
03

SGX must be enabled on provider's machine

Protected Files configuration:

- **Enclave manifest**
- **Running the docker image**

Graphene-ng Docker for Brass Golem configuration

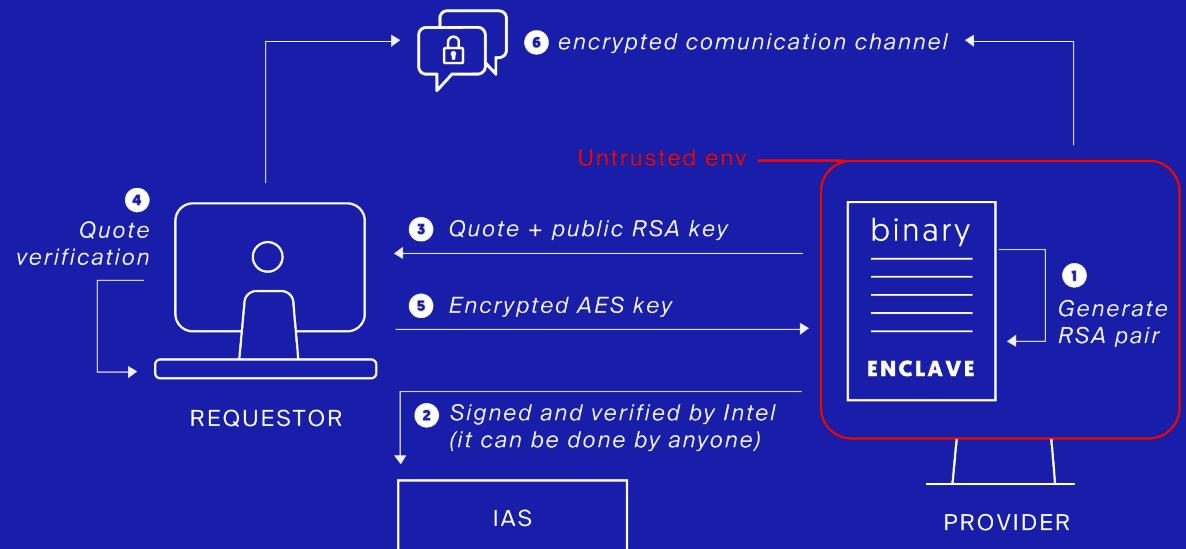


Initialization - setup and handshake

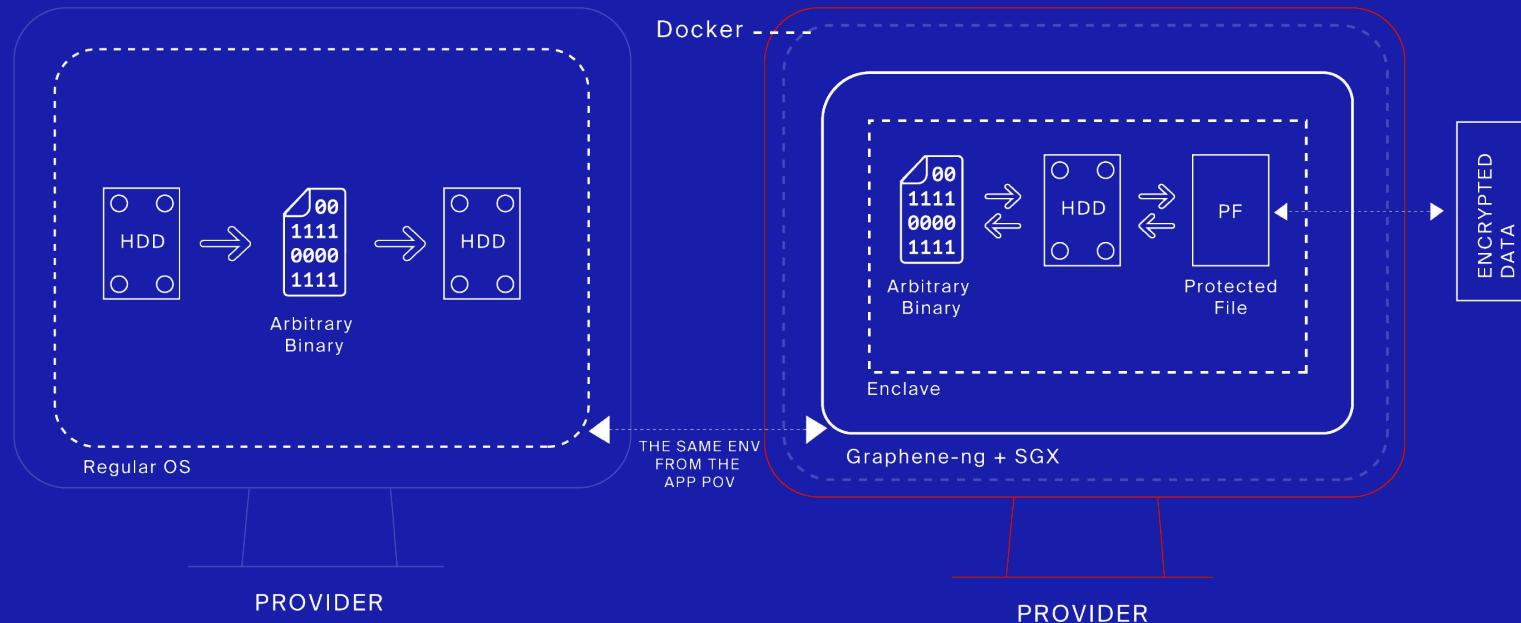
01

The process

- Semi-automatic
- Uses convenience scripts



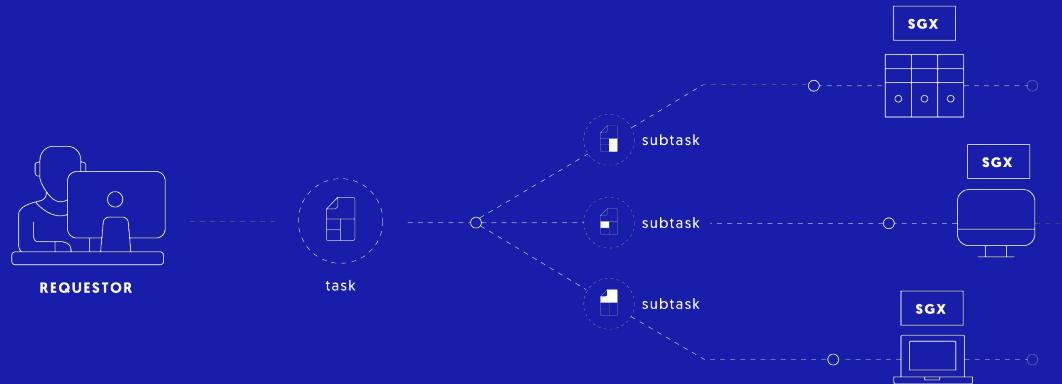
Application POV



Requestor POV

01 —————

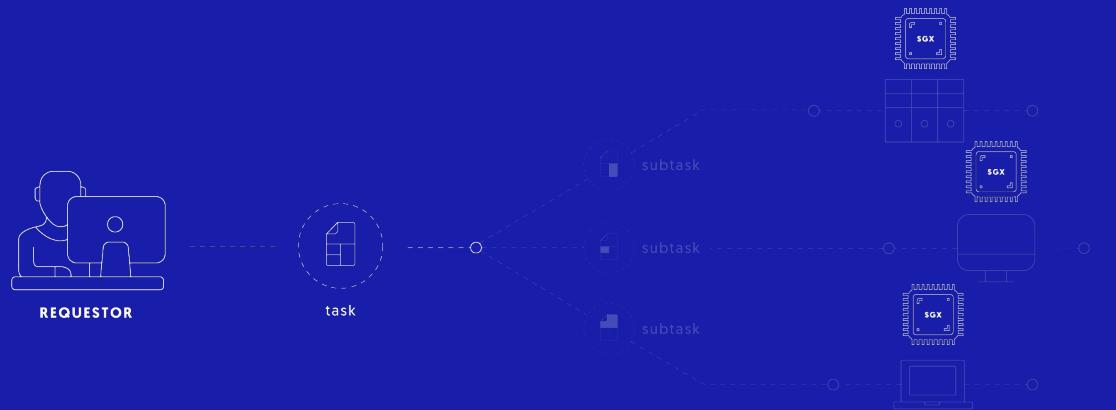
Requestor looks for SGX nodes in a P2P network



Requestor POV

02

And connects to available nodes that offer the SGX architecture



Requestor POV

03

It can be envisioned as if the local machine had more compute resources available locally

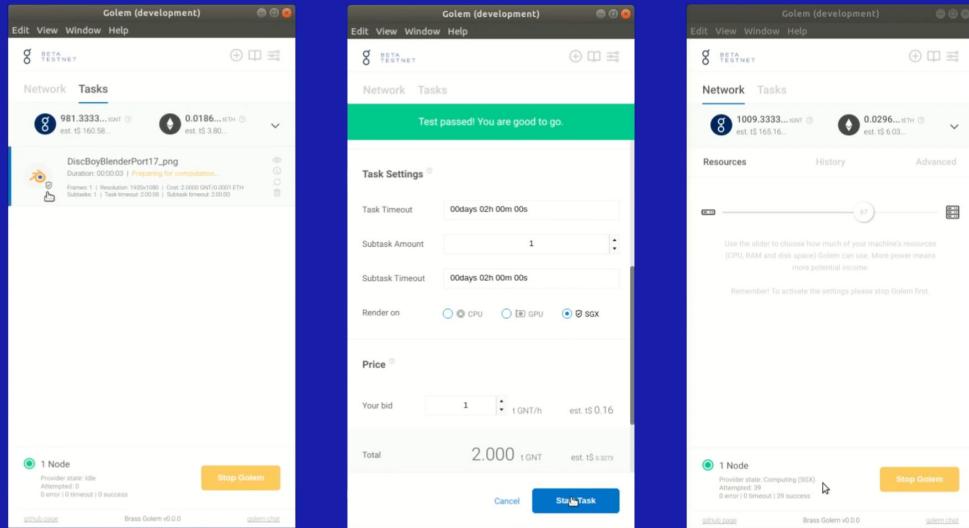


And yes, it works

01

Benefits for Golem and the community

- Confidential remote computation
- Using unmodified binaries
- SGX as an infrastructure in Golem



Next steps

01 _____

SGX:

- Liberation: IAS, FLC
- Mitigation of known attacks
- EPC size and efficiency

02 _____

Development:

- Stable and efficient Graphene-ng
- Proofs of computation
- Windows support

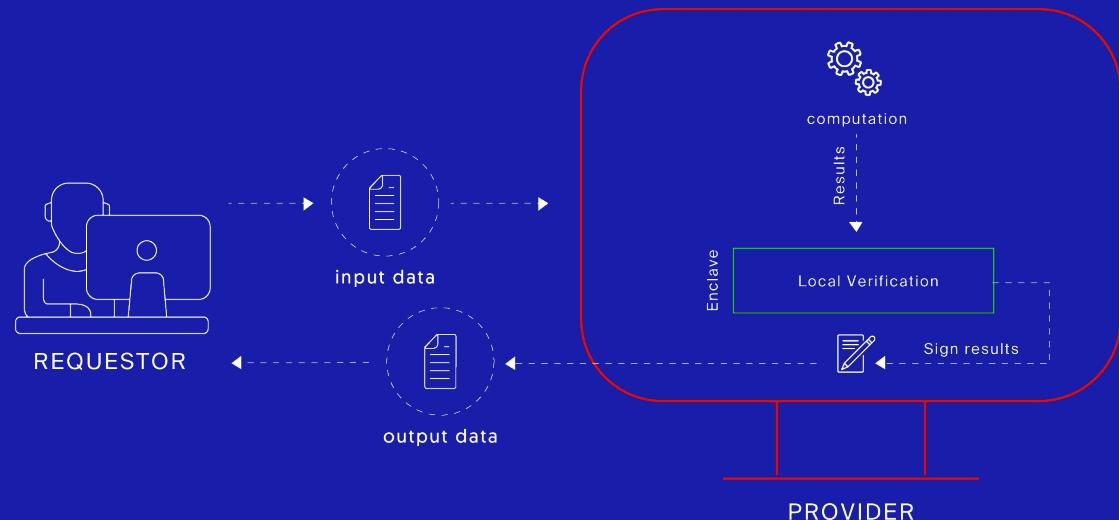
Graphene-ng Use Cases

01. Golem ecosystem

Local Verification

01

- Task is computed by an untrusted host
- And verified in an enclave run by the provider



Golem Unlimited

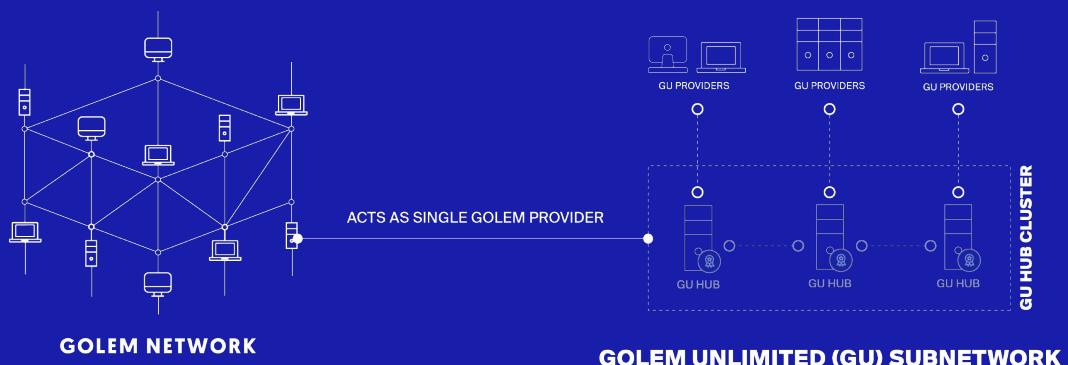
01

Golem Unlimited

- Infrastructure based on a LAN-like network setup

Two trusted actors:

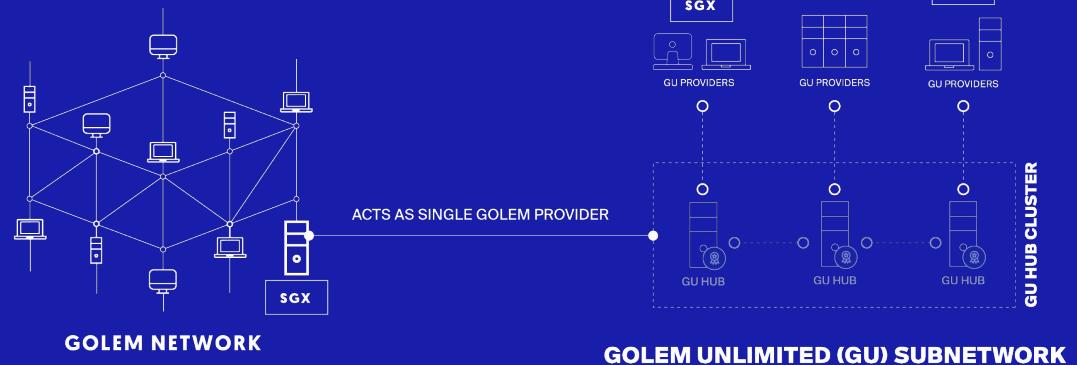
- The hub - requestor role
- Provider nodes



Golem Unlimited - powerful component with SGX

02

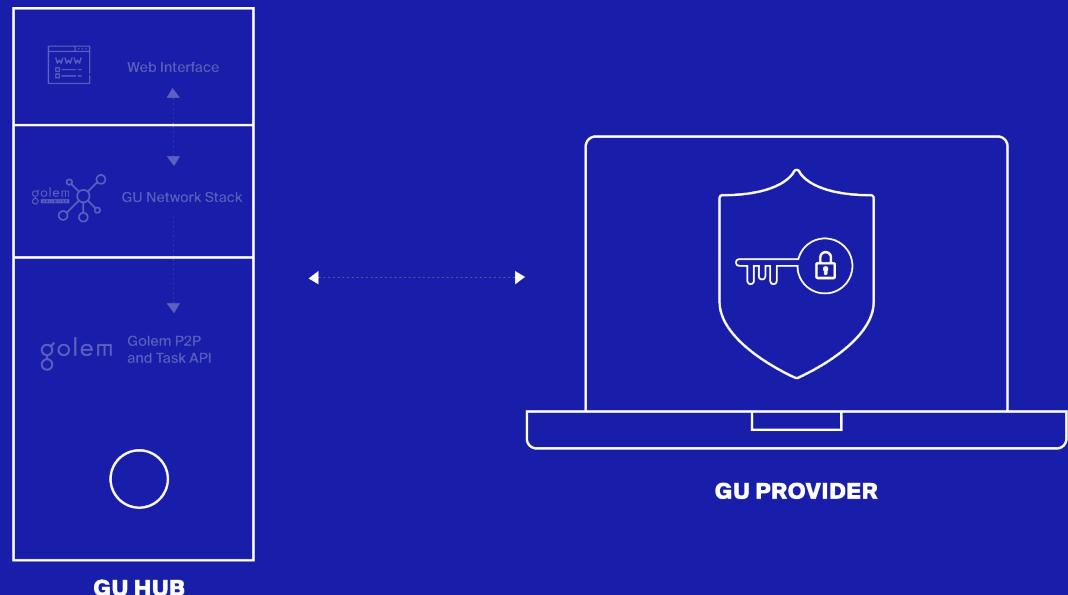
- Each Unlimited provider exposes SGX backend
- Can be envisioned as a powerful component with SGX added to the Golem network



Golem Unlimited - storing identities

03

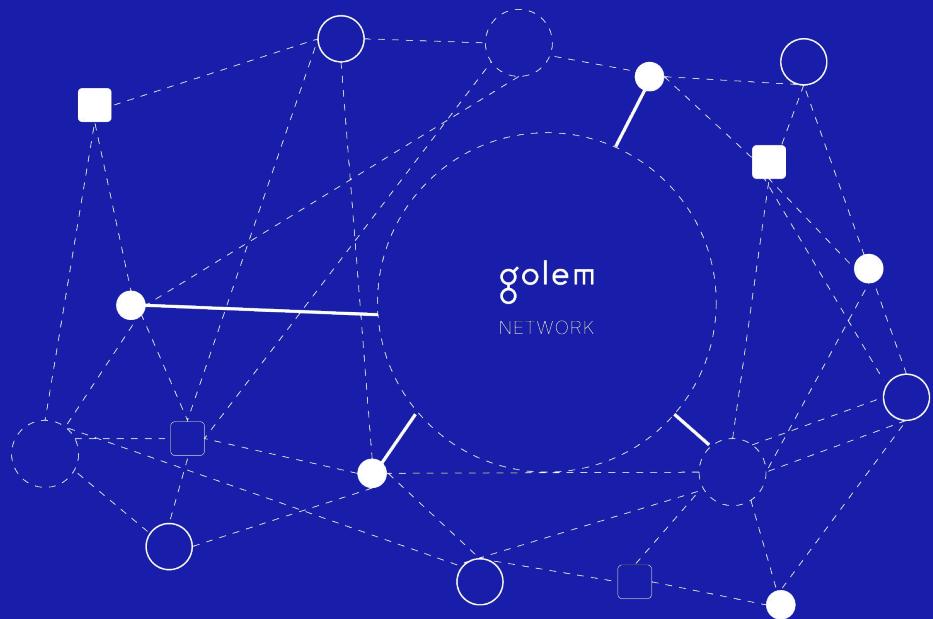
- Storing the identity information only inside enclaves
- Binds identities to physical machines in a passwordless manner



Use Cases - other Golem integrations

01

Integrations with Golem by external
developers



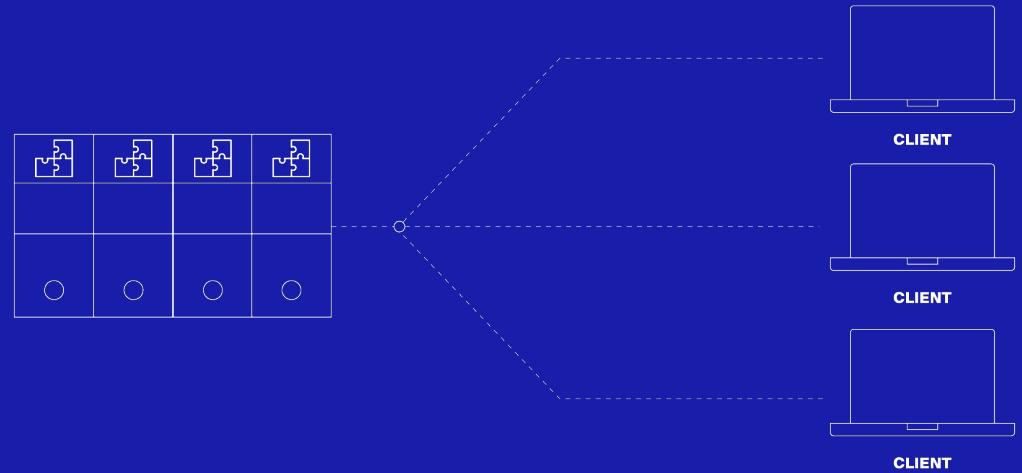
Graphene-`ng` Use Cases

02. Wider audience

Decentralized server implementation

01

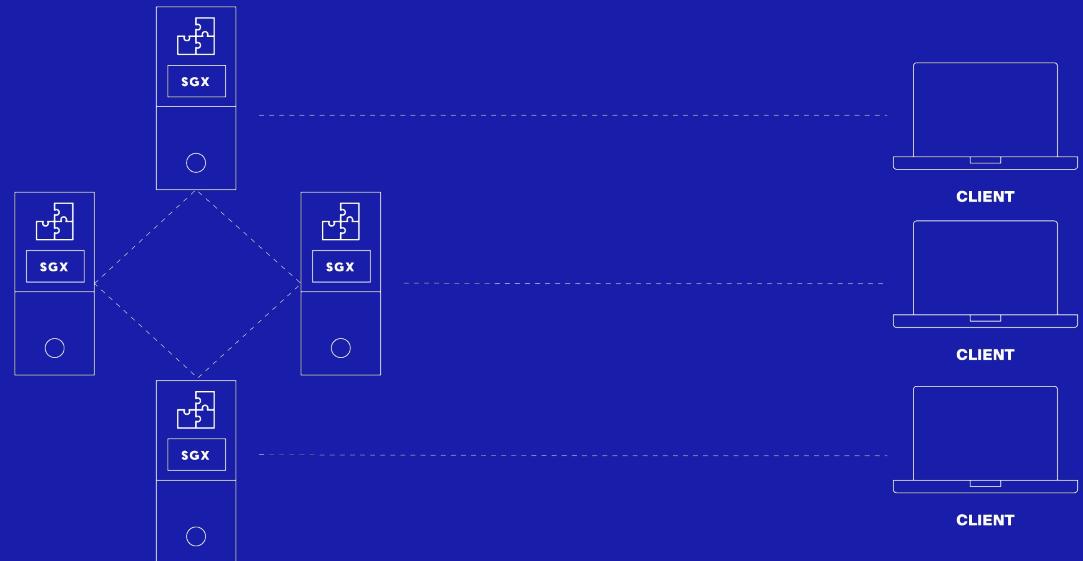
- Centralized infrastructure



Decentralized server implementation

02

- Decentralized architecture
- May be potentially used to facilitate secure multi-party computation



Atomic swap

01

Cross-blockchain atomic swap

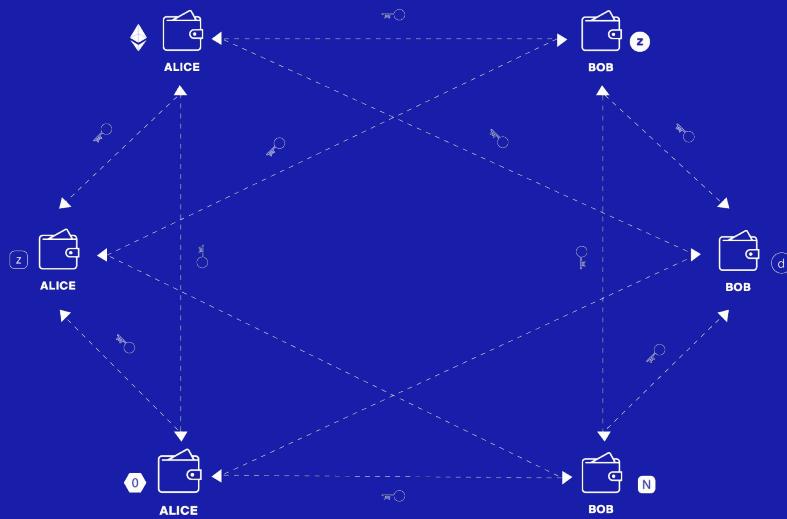
Logic implemented inside Enclave



Distributed exchanges

02

Implemented by means of cross-blockchain atomic swap



Storing secrets

01

- Sealing private keys locally
 - Hardware wallet primitive
- Sealing private keys remotely

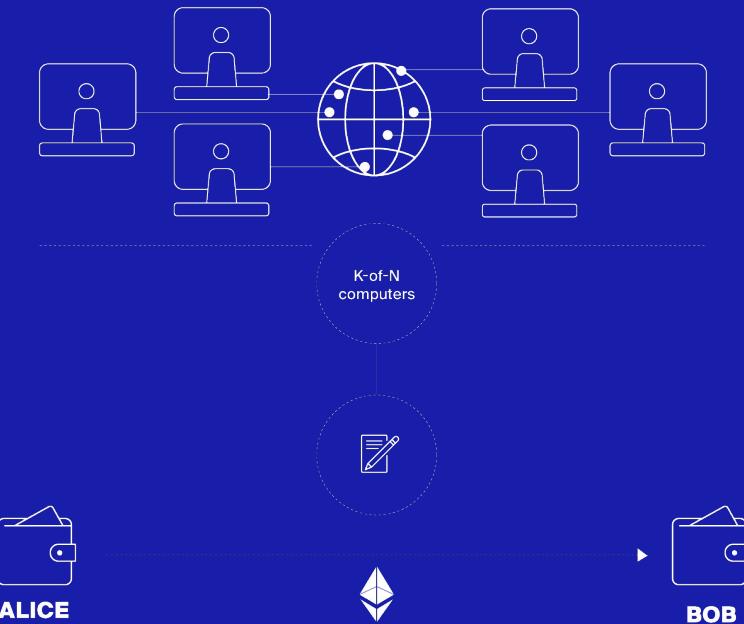


STORING PRIVATE KEYS IN CLOUD

Storing secrets

02

- Threshold signatures
- Identity stored in a decentralized way



Graphene-ng Use Cases

03. In existing projects

Minimal Viable Plasma

01

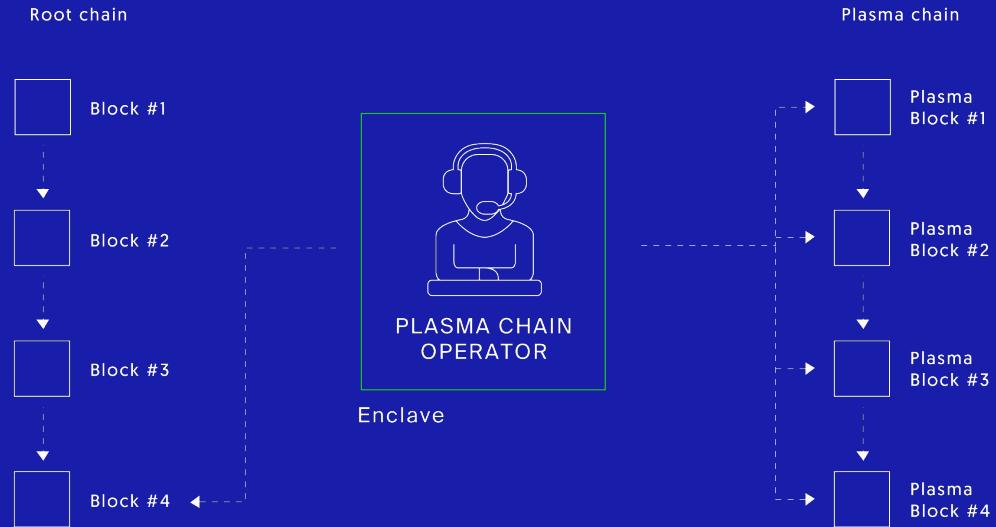
Minimal Viable Plasma with operator



Minimal Viable Plasma

02

Plasma chain operator can be implemented in an SGX enclave

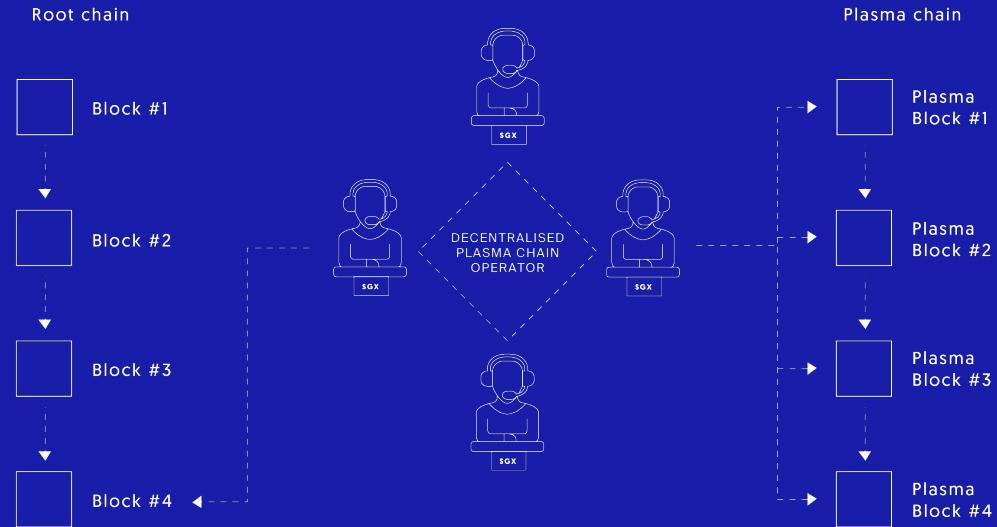


Minimal Viable Plasma

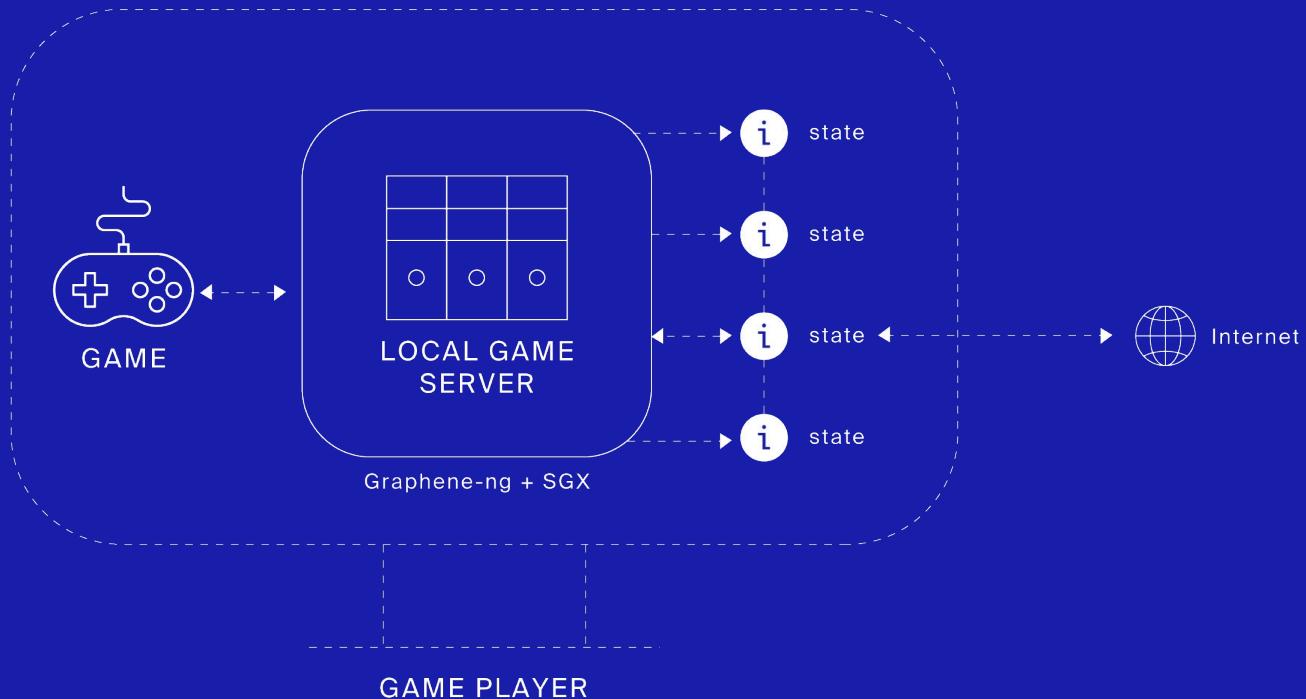
03

Plasma chain operator can be implemented in an SGX enclave

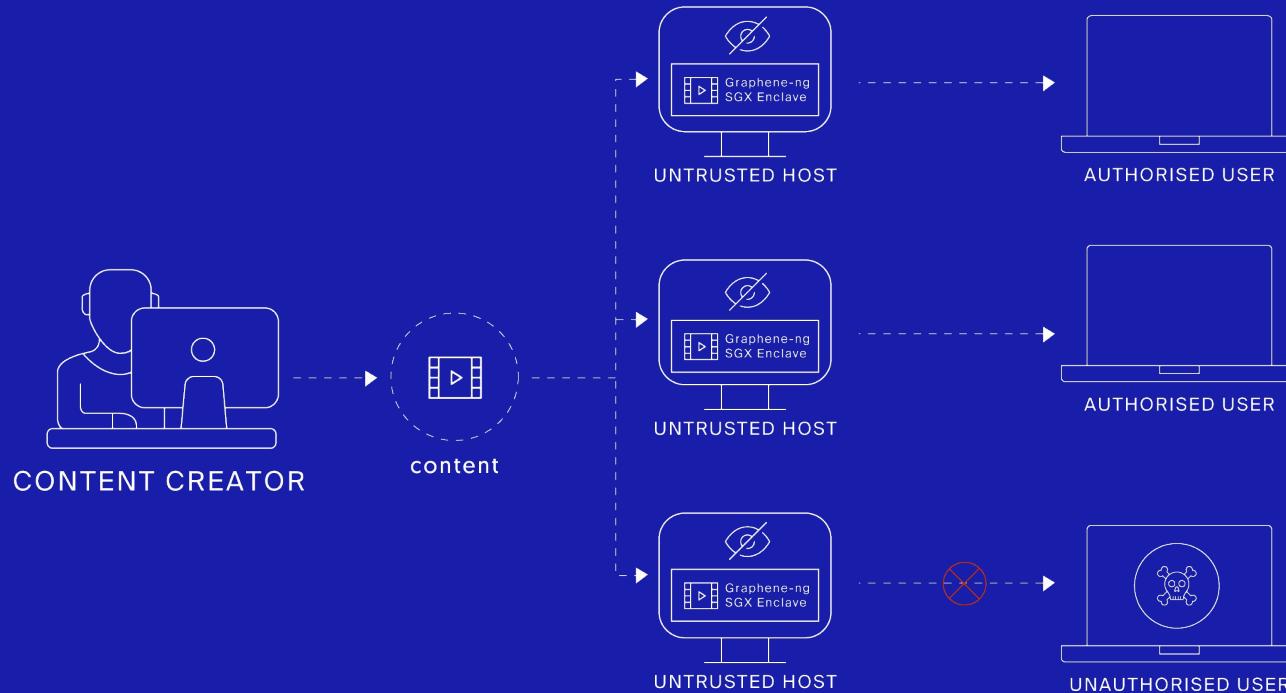
And additionally decentralised

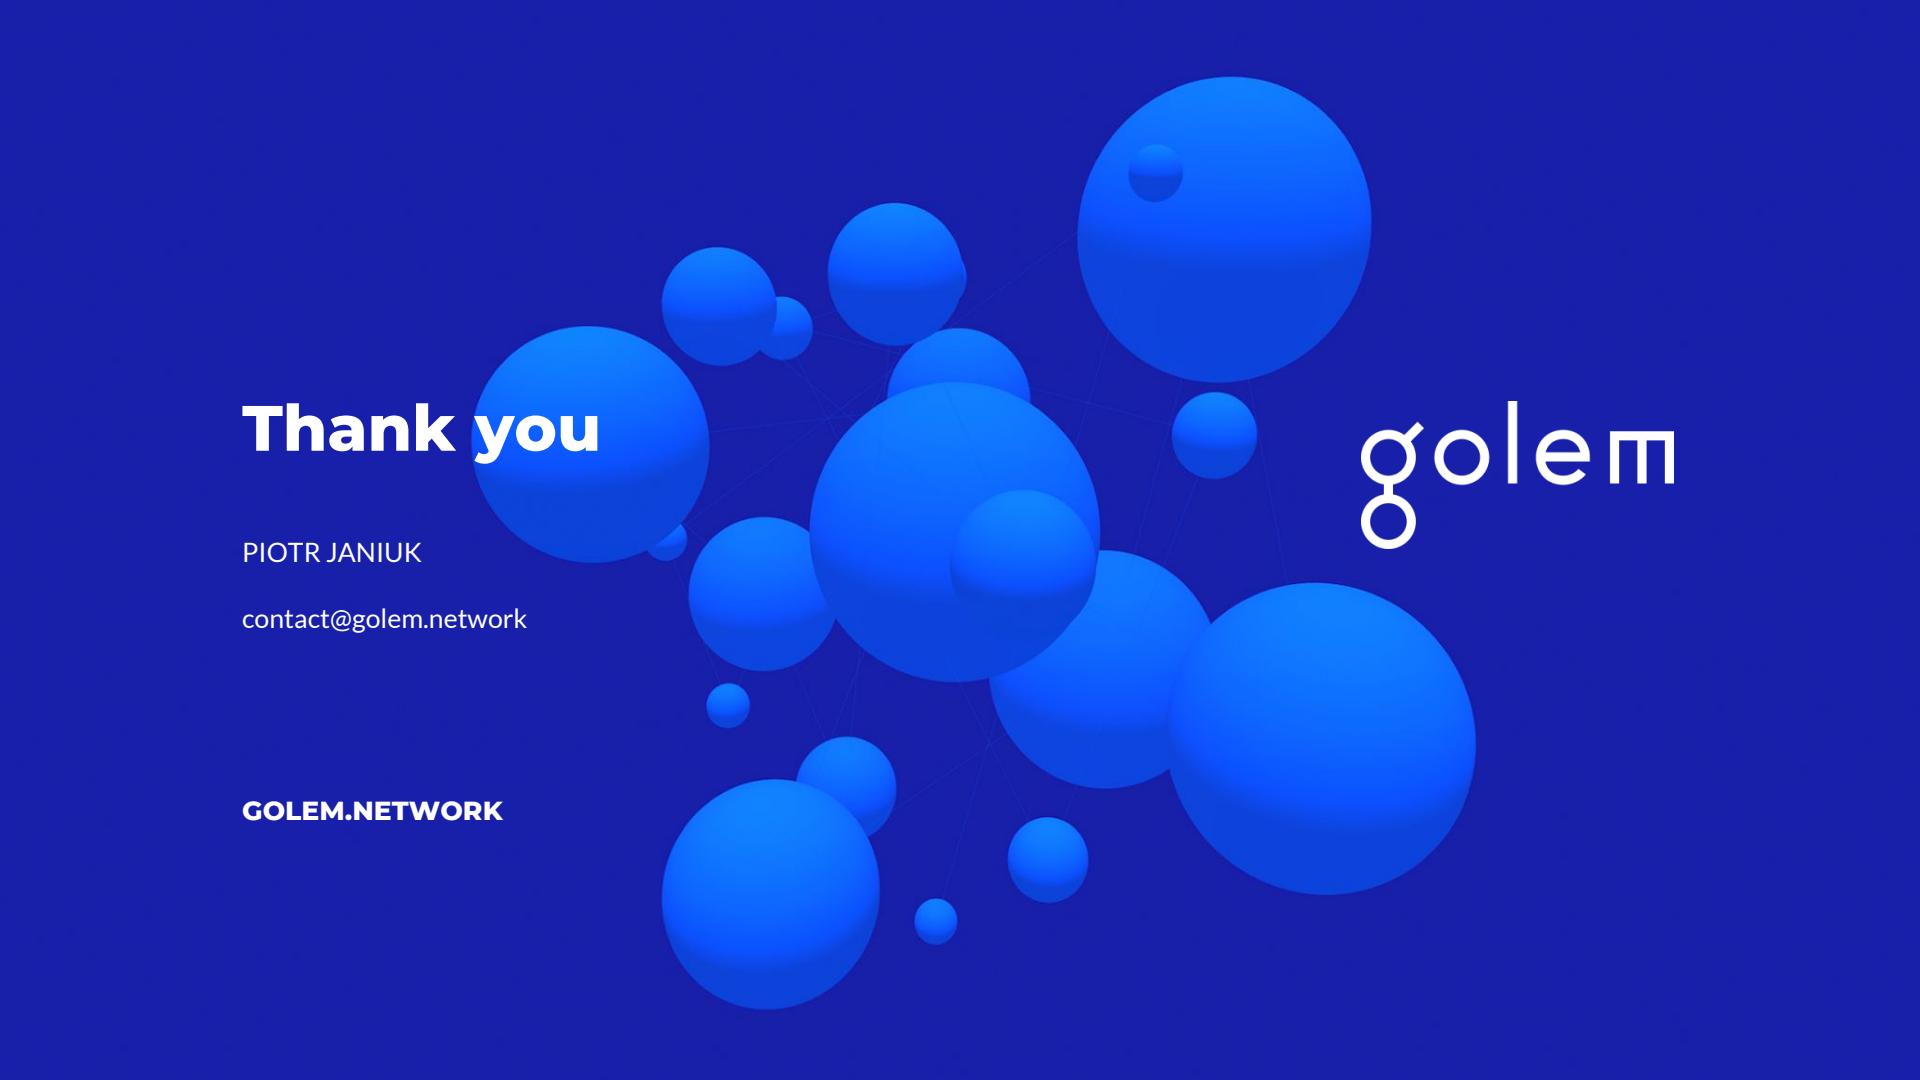


Hoard



Data streaming



The background features a dark blue gradient with a central cluster of semi-transparent blue spheres of varying sizes. A faint, light blue grid of interconnected lines forms a network pattern behind the spheres.

Thank you

PIOTR JANIUK

contact@golem.network

golem

GOLEM.NETWORK